

Monitoring and Observability

Introduction to Monitoring and Observability:


- **Observability** and **Monitoring** are often confused but are subtly different concepts, both crucial for identifying and fixing issues in software applications.
- **Goal:** To detect bugs early, ensure high-quality service, and improve user experience.

1. What is Monitoring?

- **Definition:** Monitoring involves tracking the health and performance of an application, detecting issues, and providing alerts before problems escalate.
- **Example:**
 - In an e-commerce web application, users might face latency when searching for products due to a poorly performing query in the search service.
 - Monitoring tools (e.g., Prometheus, Grafana) track metrics like **response time** and alert when a service is running slow.

Monitoring:

4 Golden Signals



- ✓ Earlier is better
- ✓ Know what to monitor
- ✓ Limit alerts

- Time it takes for a request to travel from the client to the server and back
- Number of requests a system receives over a specific period
- Percentage of requests resulting in errors, such as 404 Page Not Found or 500 Internal Server errors
- Measures resource utilization, including CPU, memory, and disk space


2. What is Observability?

- **Definition:** Observability allows for deep investigation into the root cause of issues. It provides the data (logs, metrics, and traces) necessary to troubleshoot and resolve problems.

Observability:

3 Pillars of Observability

- ✓ Decide what to log
- ✓ Cleanup logs




Provide a chronological record of events, or transactions within a system

Quantitative measurements that offer a snapshot of a system's performance over time

Helps track the flow of requests through various services and components of a system.

- **Tools for Observability:** Elasticsearch, Prometheus, Zipkin, Tempo.

Differences:



Analogy:

- **Monitoring** is like detecting an increased heart rate in a patient via vital signs.
- **Observability** is like examining the patient's full history, activities, and symptoms to find the exact cause (e.g., an allergic reaction).

Conclusion

- **Monitoring:** Detects the problem (e.g., high response time).
- **Observability:** Provides the context to identify the root cause (e.g., which microservice is the bottleneck).
- **Together:** They help teams ensure system health, user experience, and scalability.

Summary

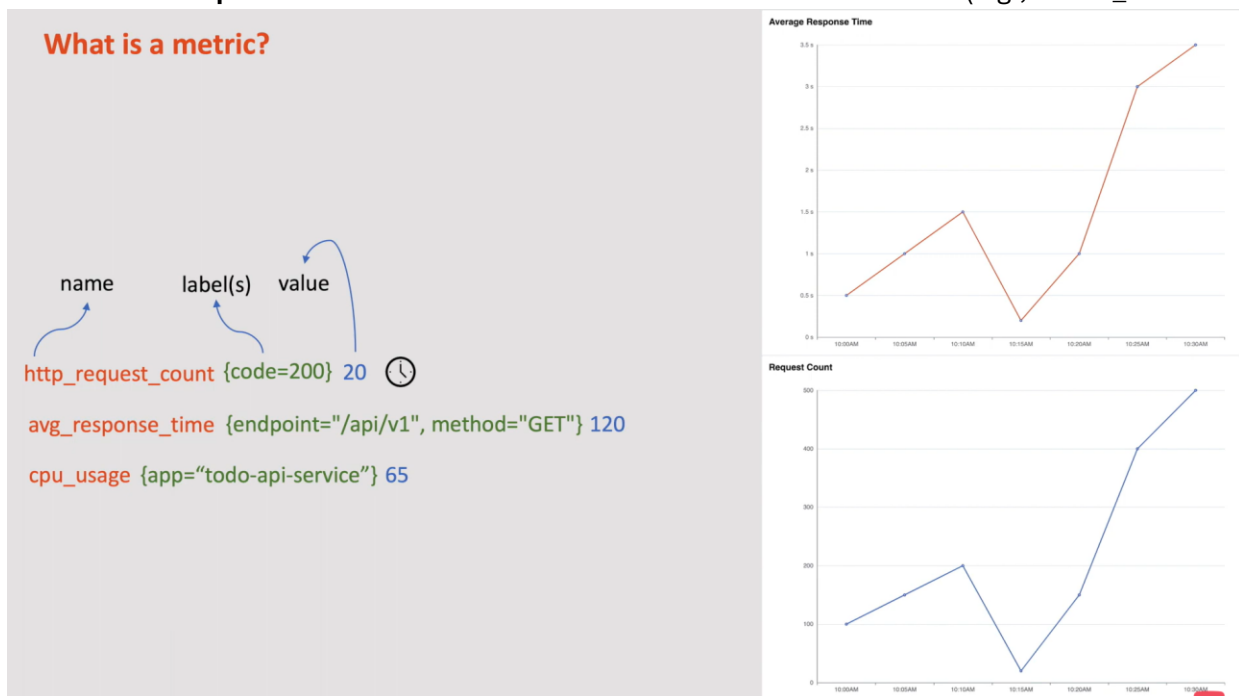
- **Monitoring** alerts you when something goes wrong.
- **Observability** helps you figure out **why** it went wrong by providing a complete set of data (logs, metrics, and traces).

Prometheus Overview:

- Prometheus is widely used for storing time series data and providing insights into system performance.

1. What is a Metric?


- **Definition:** A metric is a numerical measurement that tracks some aspect of a system's performance.
- **Example:**
 - A **request count** metric in a to-do application could indicate the number of requests made in a given time frame.
 - A metric consists of:
 - **Name:** e.g., `http_request_count`
 - **Value:** e.g., 20 requests
 - **Optional Labels:** Additional attributes to describe the metric (e.g., `status_code=200`).



2. What is a Time Series Database (TSDB)?

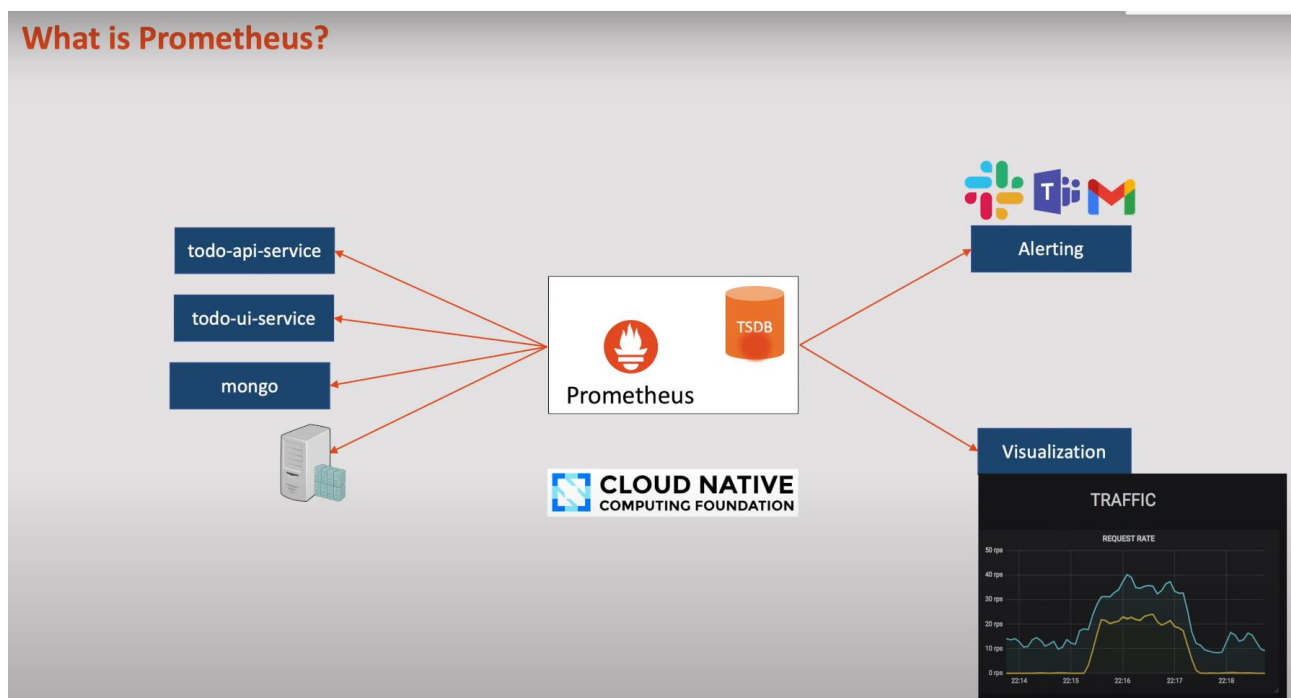
- **Definition:** A TSDB is a specialized database optimized for storing and querying time-stamped data.
- **Why TSDB?:** Traditional relational databases are not efficient for storing time series data.

- **Features:** TSDBs store data points with timestamps and associated values (e.g., request count at specific time intervals).



3. What is Prometheus?

- **Definition:** Prometheus is an open-source **observability** tool designed for collecting and storing time series data.
- **Use Cases:**
 - **Metrics Collection:** Prometheus gathers metrics from various services (e.g., applications, nodes).
 - **Centralized Metrics Storage:** Prometheus stores all metrics in a central location for easy analysis.
 - **Insights and Dashboards:** Use the data to create dashboards for monitoring system performance and behavior.
 - **Alerts:** Set up alerts for critical events like service outages or high resource usage.
- **Origin:** Developed by former Google engineers at **SoundCloud** in 2012 as an internal tool.
- **Cloud Native:** Now maintained by **CNCF** and integrated with the **cloud-native ecosystem**.
- **Native Kubernetes Support:** Prometheus provides native support for monitoring Kubernetes environments.



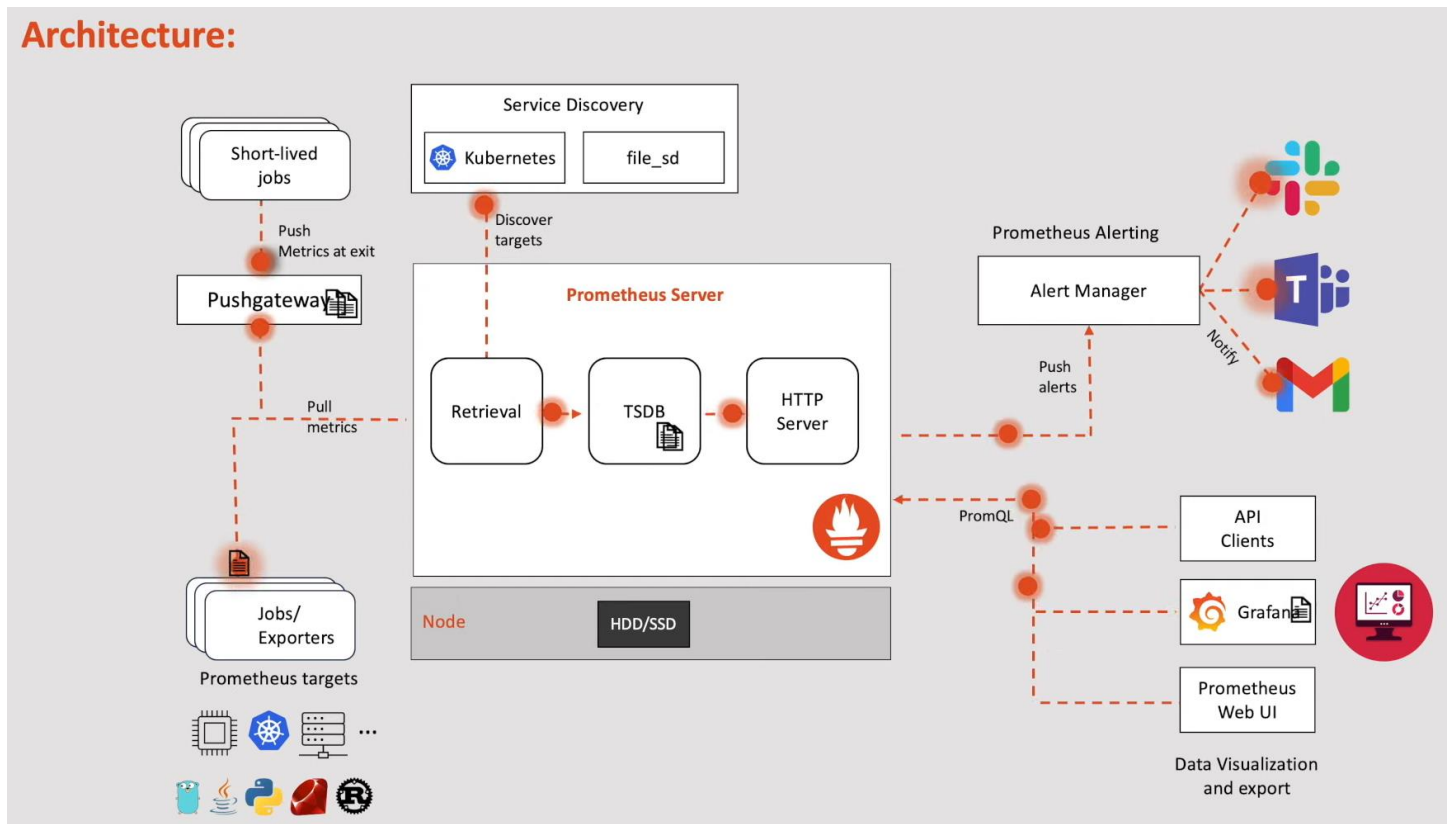
4. Facts About Prometheus

- **Open-Source:** Prometheus is an open-source tool widely adopted due to its community-driven nature.
- **Integration with Kubernetes:**
 - Prometheus easily integrates with Kubernetes, as many Kubernetes components expose metrics in the Prometheus format.
- **Comparison with Other Tools:**
 - While there are other monitoring solutions (e.g., Nagios, DataDog, New Relic), Prometheus is the **default standard** for Kubernetes monitoring.

Summary

- **Prometheus** is a powerful, open-source tool that enables the collection, storage, and analysis of time series data for observability and monitoring.
- It is widely used for tracking system metrics, setting up alerts, and gaining insights into application performance.

Architecture:



1. Prometheus Server

- **Core Component:** The **Prometheus Server** is the central part responsible for collecting and storing metrics.
- **Metrics Collection:** Prometheus collects metrics by scraping **HTTP endpoints** from various targets, such as:
 - Applications
 - Bare metal servers
 - Kubernetes clusters
 - Database instances

- **Time Series Database (TSDB):**
 - Metrics are stored in a **local time series database**.
 - Default retention is **15 days**, but this can be configured or stored remotely (e.g., on S3).
- **API:** Prometheus provides an **API** to query the stored data, which can be accessed via **Grafana** or the **Prometheus UI**.

2. Client Libraries

- **Purpose:** Custom metrics for applications can be generated using **client libraries**.
- **Supported Languages:** Prometheus supports client libraries for various programming languages, such as **Go, Java, Python, Ruby, and Rust**.
- **Instrumentation:** Adding code to an application to expose its own metrics is called **instrumentation**. For example, tracking the number of to-do items in a to-do application.

3. Exporters

- **Purpose:** For systems that cannot be instrumented directly (like **Linux** or **MongoDB**), **exporters** help fetch existing metrics and expose them in Prometheus-compatible formats.
- **How It Works:**
 - Exporters fetch metrics from systems and expose them via HTTP endpoints.
 - Prometheus scrapes these endpoints to collect the data.
- **Examples of Exporters:**
 - **Node Exporter:** Fetches metrics from Linux systems.
 - **MongoDB Exporter:** Fetches metrics from MongoDB.

4. Push Gateway

- **Use Case:** For **short-lived applications** like batch jobs or Lambda functions that run briefly and exit, Prometheus's **pull-based system** doesn't work well.
- **Solution:** The **Push Gateway** allows these applications to **push** their metrics, and Prometheus can scrape the data from the Push Gateway.
- **Important Note:** The Push Gateway can introduce additional storage requirements, so it should be used mindfully.

5. Service Discovery

- **Purpose:** In a **containerized environment** (e.g., Kubernetes), targets like **pods** are frequently added and removed, making manual configuration of Prometheus for each new target cumbersome.
- **Solution:** Prometheus has **service discovery** that automatically detects and scrapes targets as they are added or removed.
- **Benefit:** Automated monitoring and scraping without manual intervention, which is critical in dynamic environments like Kubernetes.

6. PromQL and Prometheus UI

- **PromQL: Prometheus Query Language (PromQL)** is used to query and analyze the stored metrics.
- **Prometheus UI:**
 - **Access:** By default, accessible on port **1990** of the Prometheus server.
 - **Purpose:** Useful for ad-hoc queries and debugging.
 - **Limitations:** Not ideal for creating dashboards. For visualization, **Grafana** is recommended.

7. Alert Manager

- **Purpose:** Prometheus can trigger alerts based on specific conditions or thresholds.
- **How It Works:**
 - **Alert Rules:** Set conditions for alerts.
 - **Alert Manager:** When an alert condition is met, the Alert Manager sends notifications through various channels, such as **Slack, Email, or Microsoft Teams**.
- **Role:** The **Alert Manager** ensures Prometheus is also used for **alerting** alongside monitoring.

Summary

- **Prometheus** is an all-in-one tool for monitoring and alerting:
 - Collects metrics from configured targets at regular intervals.
 - Stores metrics in a time series database (TSDB).
 - Provides a UI and supports PromQL for querying metrics.
 - Sends notifications via the Alert Manager when defined thresholds are met.
- **Components:** Key components include the Prometheus Server, Client Libraries, Exporters, Push Gateway, Service Discovery, PromQL, and Alert Manager. Together, these components enable effective monitoring and alerting.

Prometheus Installation:

Installation using helm -> https://youtu.be/kmjfRm82Sms?si=KycIB2AitLmeo_f0